

# Kubernetes

---



# Kubernetes Agenda

---

- A bit about Kubernetes
- Terminology
- K8S in depth
- Demos (For most sections)
- Hands on

# What is Kubernetes?

---

- The meaning of **Kubernetes** is **Captain** in Greek
- **Kubernetes** is also referred to as K8S (**K**ubernetes S)
- Open-source container-orchestration software
- Supply automation for deployment, scaling, operation and more
- **Kubernetes** is a platform which manage the containers for you (deploy health-check etc)
- Originally Developed by Google (2013) and v1.0 released to the public in 07/2015
- Was based upon Google **Borg** (Large Cluster Management) developed internally by Google
- Today is managed by the Cloud Native Computing Foundation (CNCF)
- The most contributed project by Unix community after the Linux Kernel

# What is Kubernetes? [batteries-included]

- Think of Kubernetes like a higher-level language for your application architecture.
  - You describe what you need and it tries to resolve it
  - K8S acts as an engine to **resolve state** using the **abstracted resources** to **deploy and manage** your application.
- It is **declarative**, and not **imperative**.
- You tell it or “declare” what you want, and it figures out the rest.

imperative	you tell it what you want it <b>in every step</b>
declarative	you tell it what you want it to be, and <b>it figures it out</b>

- K8S gives you: **self-healing** and **seamless upgrading or rollback** of applications

# What can Kubernetes REALLY do?

---

- Autoscale Workloads
- Blue/Green Deployments
- Manage Stateless and Stateful Applications
- Provide native methods of service discovery (replaceable)
- Easily integrate and support 3rd party apps (ex Helm, containers)
- **The killing feature:**

Use the SAME API across everywhere .... (bare metal & cloud providers)

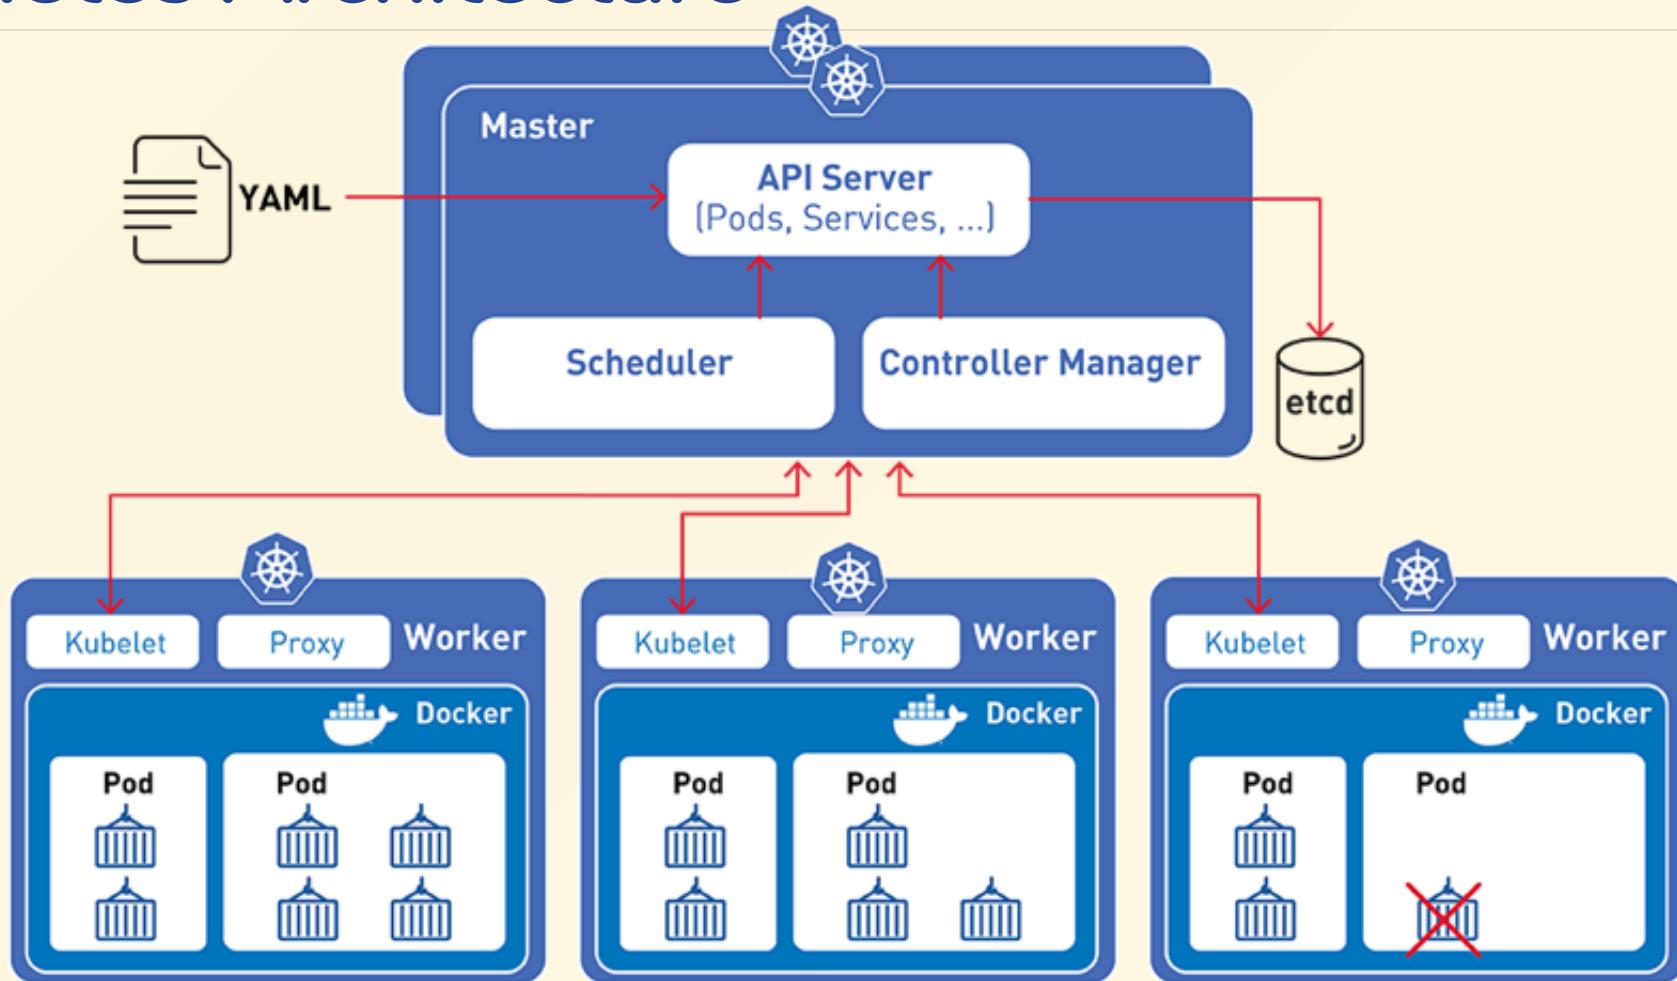
---

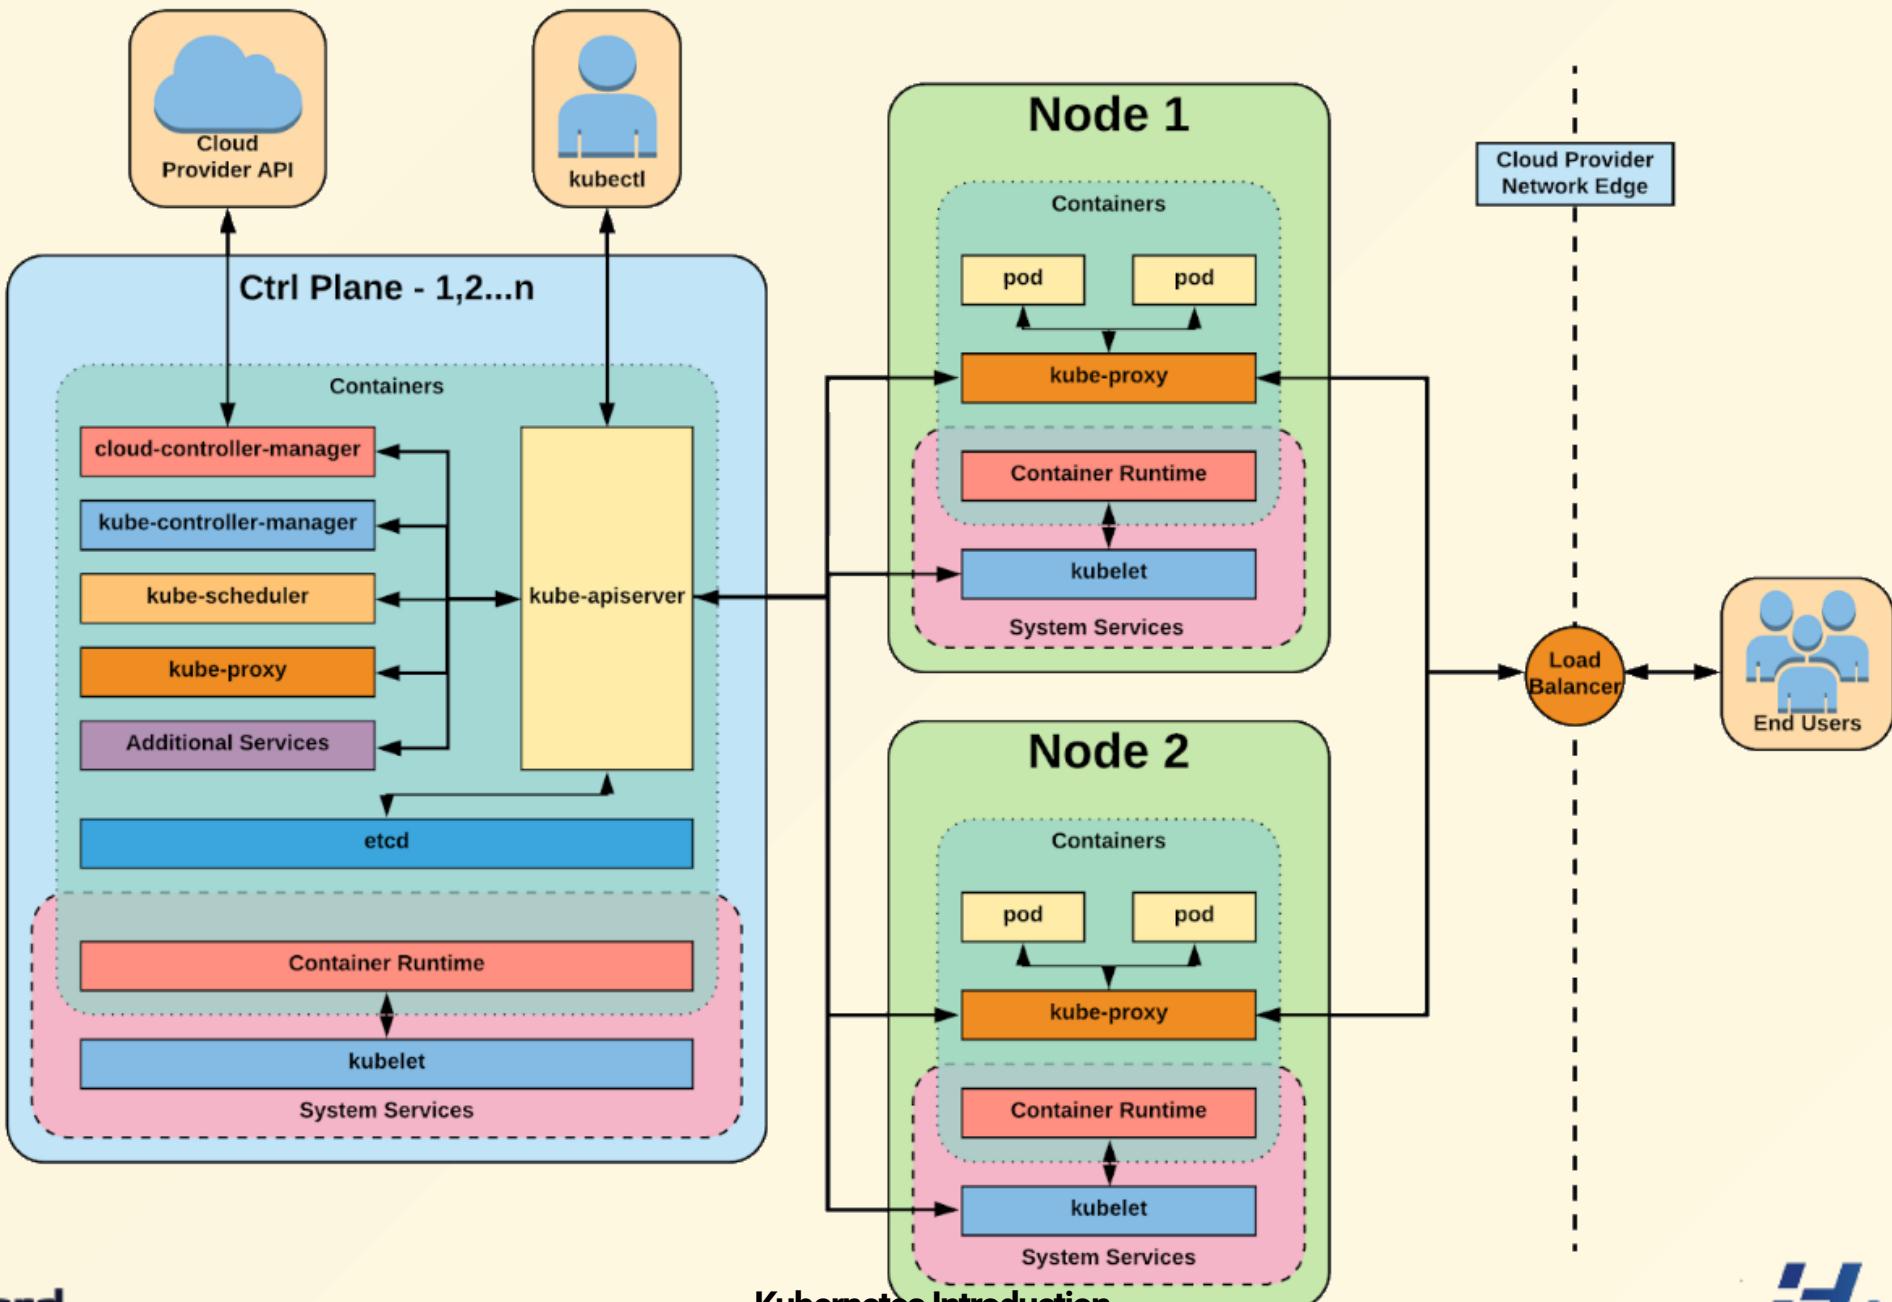
# Some Numbers

---

- Over 61,000 stars on Github
- ~2400 Contributors to K8s Core
- Most discussed repo on github
- 1000+ pull requests
- ~90,000 commits !!!!

# Kubernetes Architecture





## Kubernetes Introduction

© CodeWizard ltd | [nirgeier@gmail.com](mailto:nirgeier@gmail.com)

# Kubernetes resources

- The Kubernetes API defines a set of resources
- Resources are organized by type, `kind`  
Some common resource `kinds` are:

Kind	Description
<code>node</code>	a machine — physical or virtual — in our cluster
<code>pod</code>	group of containers running together on a host
<code>service</code>	stable network endpoint to connect to one or multiple containers
<code>namespace</code>	isolated group of things
<code>replicaset</code>	set of containers which can be scaled
<code>secret</code>	sensitive data to be passed to a container

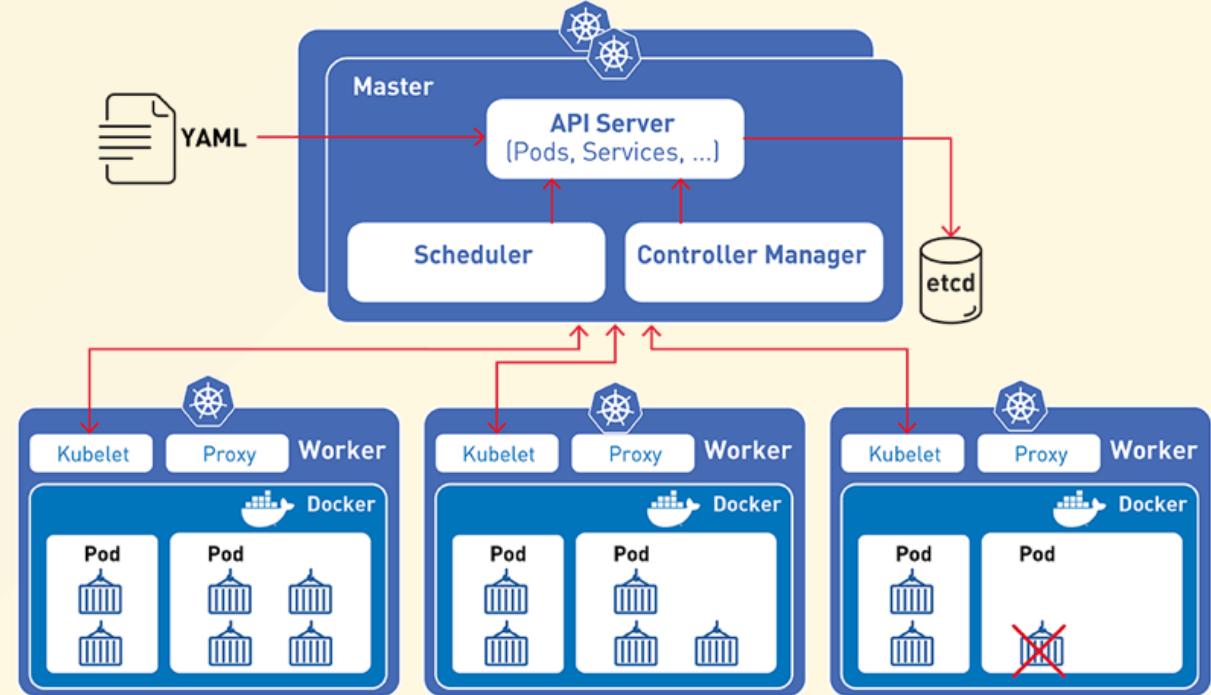
# Terminology / Main features

---

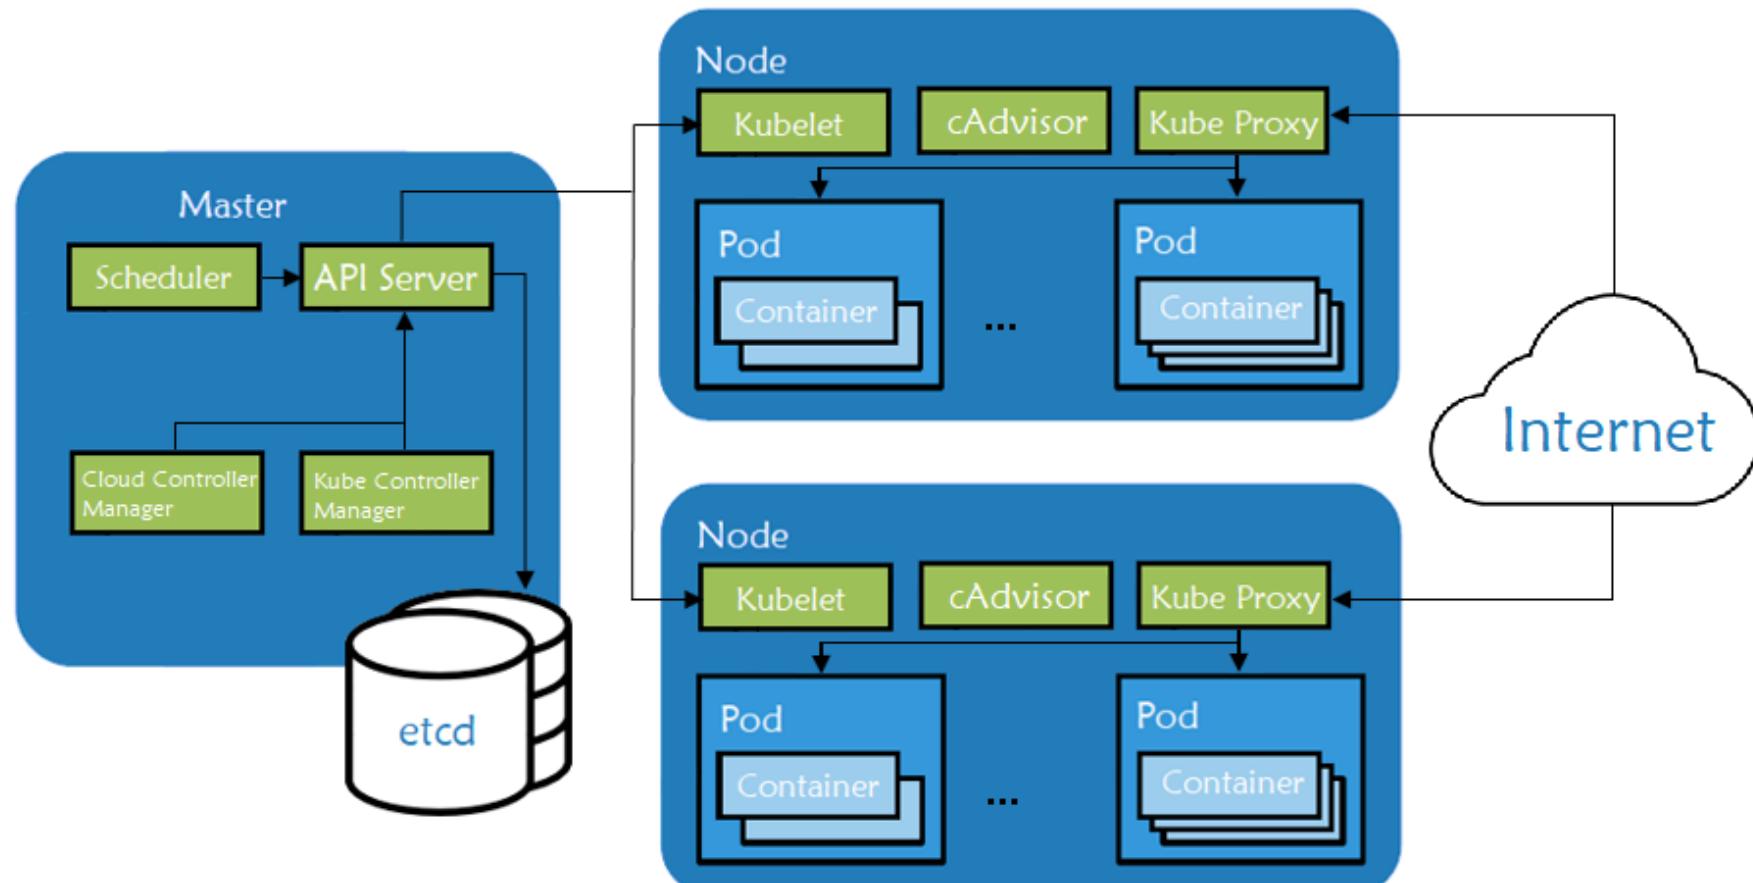
- Cluster
- Nodes / Master node
- Pod
- Replica Sets / Deployments
- Services
- Volumes
- Namespaces
- ConfigMaps and Secrets
- Labels and selectors
- etcd

# Terminology - Cluster

- When you deploy Kubernetes, you get a **cluster**.
- A **cluster** is a **set of machines**, called **nodes**, that run containerized applications managed by Kubernetes



# Terminology - Nodes (architecture)



# Terminology - Nodes

---

- A node is a **worker machine** in Kubernetes, previously known as a **minion**.
- A node may be a VM or physical machine, depending on the cluster
- The services on each node include the container runtime, kubelet and kube-proxy
- k8s has a unique node named **Master Node** (Master Components)
- Node is a top-level resource in the Kubernetes REST API
- Each node contains the services necessary to run pods and is managed by the master components
  - Container Engine (typically Docker)
  - **kubelet** (aka “node agent”)
  - **kube-proxy** (a necessary but not sufficient network component)

# Terminology - Master Node (Master Components)

---

- Master components in Kubernetes is a collection of services
- Master components provide the cluster's control plane
- Master components manage the cluster (ex: scheduling, deploying)
- Master components can run on any node but usually deployed to a dedicated node
- There can be multiple master nodes for high availability
- The main components of the **Master Node** (in addition to worker node components) are:
  - kube-apiserver
  - etcd (a highly available key/value store; the “database” of Kubernetes)
  - kube-scheduler
  - kube-controller-manager
  - Node Controller

# Terminology - Pod

---

- The most basic component of k8s
- Pod is a **group of containers** (one or more) which is **deployed** to the host machine
- Pod describe the application running on k8s
- **Pods serve as unit of deployment for horizontal scaling, and replication**
- Each pod get its unique ip in the cluster
- Pods can communicate with each other using the Pod IP address,  
Developer should use Services and not Pod IP
- Pods can share storage (volumes), network, host resources (cpu, memory)
- Pods are managed through the **k8s API**
- Pods run in a shared context (Linux namespaces, cgroups)
- Applications with in the same pod can use *localhost* for communication

# Terminology - Pod (sample - yaml file)

- kind: pod

```
apiVersion: v1
kind: Pod <-----<<
metadata:
  name: pod1
  labels:
    tier: frontend
spec:
  containers:
  - name: hello1
    image: gcr.io/google-samples/hello-app:2.0
    imagePullPolicy: Always
    command: ["echo", "SUCCESS"]
```

```
# pull the image and create a container
$ kubectl create -f <file name>
```

# Basic pod commands

```
# Get k8s kinds (resources) like pods, deployments, replicateSets  
kubectl get <...>  
  
# Describe the given pod  
kubectl describe pod <pod>  
  
# Create new service (= expose new port)  
kubectl expose pod <pod> --port=<port> --name=<name>  
  
# Port forward from the pod to the host  
kubectl port-forward <pod> <port>  
  
# Execute command on the pod  
kubectl exec <pod> --command ...  
  
#Execute interactive pod, run pod with shell - !!!! Usefull for debugging pods  
kubectl run -i -tty <name> --image=... --restart=Never --sh
```

# Terminology - Replica Sets

- The aim of **ReplicaSet** is to maintain a stable set of replica Pods running at any given time
- **ReplicaSet** is a collection of definitions which specify the pods.
- It contains pod templates for creating or updating new pods.
- In production its much better to use **Deployment** than **ReplicaSet**, deployment has more features
- To get the ReplicaSet information:

```
kubectl get rs
```

# Terminology - Replica Sets

```
apiVersion: apps/v1 # our API version
kind: ReplicaSet    # The kind we are creating
Metadata: # Specify all Metadata like name, labels
  name: some-name
  labels:
    app: some-App
    tier: some-Tier
Spec:
  replicas: 3 # Here is where we tell k8s how many replicas we want
  Selector: # This is our Label selector field.
    matchLabels:
      tier: some-Tier
    matchExpressions:
      - {key: tier, operator: In, values: [some-Tier]} # we are using the set-based operators
  template:
    metadata:
      labels:
        app: some-App
        tier: someTier
  Spec: # This spec section should look like spec in a pod definition
    Containers:
```

# Terminology - Deployments

---

- Deployment is the pod(s) specifications (labels, replicas, network etc)
- Deployment provides declarative updates for Pods and ReplicaSets .
- Deployment control and update the state of the Pods in the ReplicaSet
- Main Deployment use cases
  - Update ReplicaSet for rollback
  - Update Pods (add/remove/update)
  - Scale up/down
  - Clean old ReplicaSet

# Terminology - Deployments

---

- Deployment ensures that only a certain number of Pods are down while they are being created or updated.
- Every time deployment is noticing changes, a ReplicaSet is created to reflect the changes to the desired Pods
- During update ,  
by default, at least 75% of the desired number of Pods are up (25% max unavailable).
- During creation ,  
by default, at most 125% of the desired number of Pods are up (25% max surge).

# Terminology - Deployments (yaml fields)

## revisionHistoryLimit

- The number of previous iterations of the Deployment to retain.

## strategy:

- Describes the method of updating the Pods based on the type.
- Valid options are Recreate or RollingUpdate.

Recreate	All existing Pods are killed before the new ones are created.
RollingUpdate	Cycles through updating the Pods according to the parameters: maxSurge and maxUnavailable.

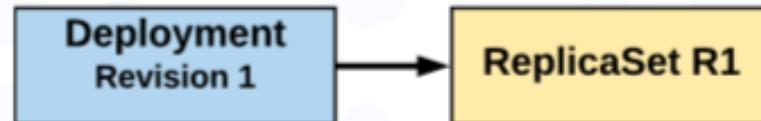
# RollingUpdate Deployment

Updating pod template generates a new **ReplicaSet** revision.

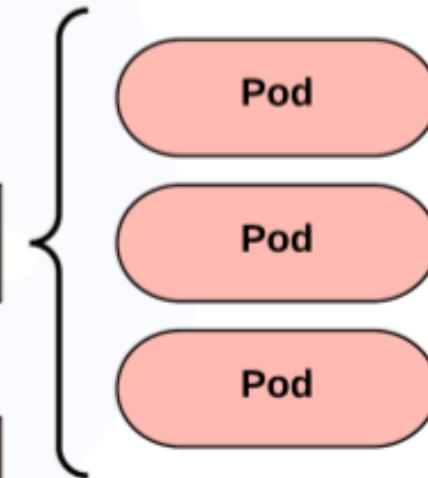
```
R1 pod-template-hash:  
676677fff  
R2 pod-template-hash:  
54f7ff7d6d
```

```
$ kubectl get replicaset  
NAME          DESIRED  CURRENT  READY   AGE  
mydep-6766777fff    3        3        3      5h
```

```
$ kubectl get pods  
NAME          READY   STATUS    RESTARTS   AGE  
mydep-6766777fff-9r2zn  1/1     Running   0          5h  
mydep-6766777fff-hsfz9  1/1     Running   0          5h  
mydep-6766777fff-sjxhf  1/1     Running   0          5h
```



ReplicaSet R2



# RollingUpdate Deployment

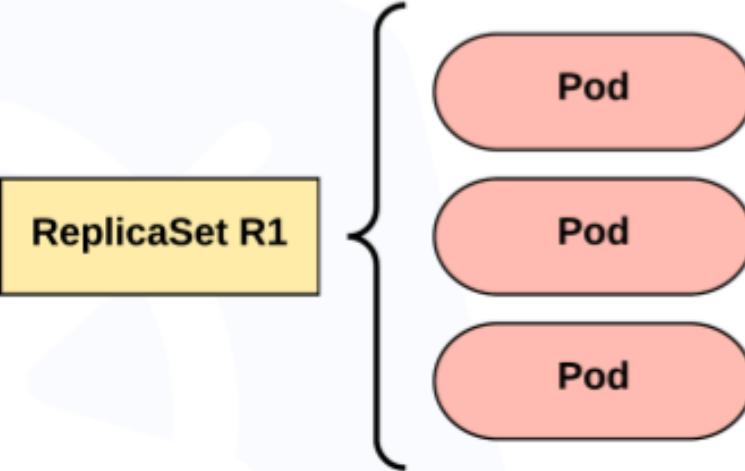
New **ReplicaSet** is initially scaled up based on [maxSurge](#).

```
R1 pod-template-hash:  
676677fff  
R2 pod-template-hash:  
54f7ff7d6d
```

```
$ kubectl get replicaset  
NAME          DESIRED  CURRENT  READY   AGE  
mydep-54f7ff7d6d  1        1        1      5s  
mydep-676677fff  2        3        3      5h
```

```
$ kubectl get pods  
NAME          READY   STATUS    RESTARTS   AGE  
mydep-54f7ff7d6d-9gv1l  1/1     Running   0          2s  
mydep-676677fff-9r2zn  1/1     Running   0          5h  
mydep-676677fff-hsfz9  1/1     Running   0          5h  
mydep-676677fff-sjxhf  1/1     Running   0          5h
```

Deployment  
Revision 2



ReplicaSet R2



# RollingUpdate Deployment

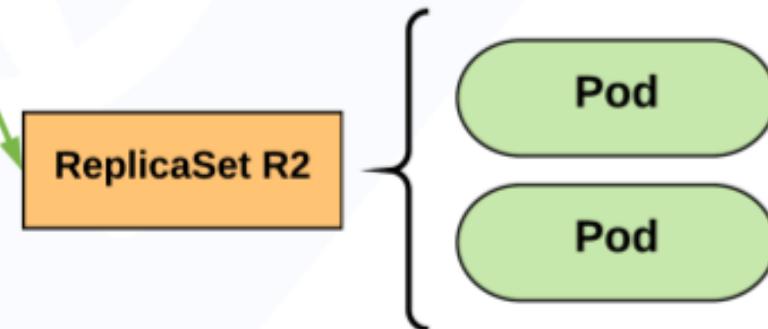
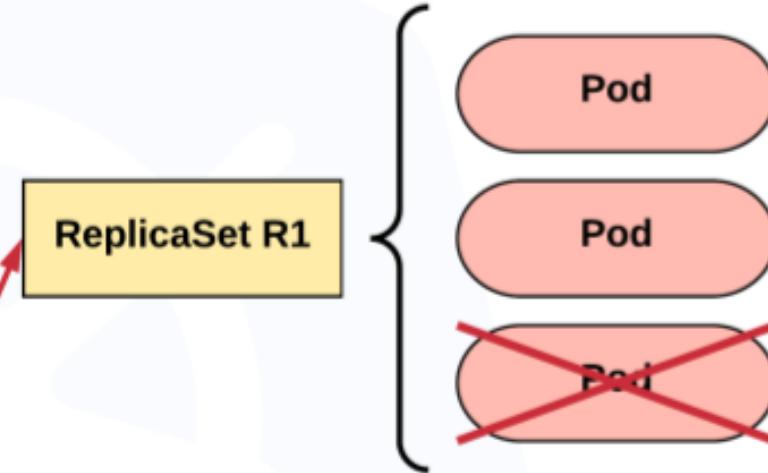
Phase out of old Pods managed by **maxSurge** and **maxUnavailable**.

```
R1 pod-template-hash:  
676677fff  
R2 pod-template-hash:  
54f7ff7d6d
```

Deployment  
Revision 2

```
$ kubectl get replicaset  
NAME          DESIRED  CURRENT  READY   AGE  
mydep-54f7ff7d6d  2        2        2      8s  
mydep-676677fff  2        2        2      5h
```

```
$ kubectl get pods  
NAME          READY   STATUS    RESTARTS   AGE  
mydep-54f7ff7d6d-9gv1l  1/1     Running   0          5s  
mydep-54f7ff7d6d-cqv1q  1/1     Running   0          2s  
mydep-676677fff-9r2zn  1/1     Running   0          5h  
mydep-676677fff-hsfz9  1/1     Running   0          5h
```



# RollingUpdate Deployment

Phase out of old Pods managed by **maxSurge** and **maxUnavailable**.

```
R1 pod-template-hash:  
676677fff  
R2 pod-template-hash:  
54f7ff7d6d
```

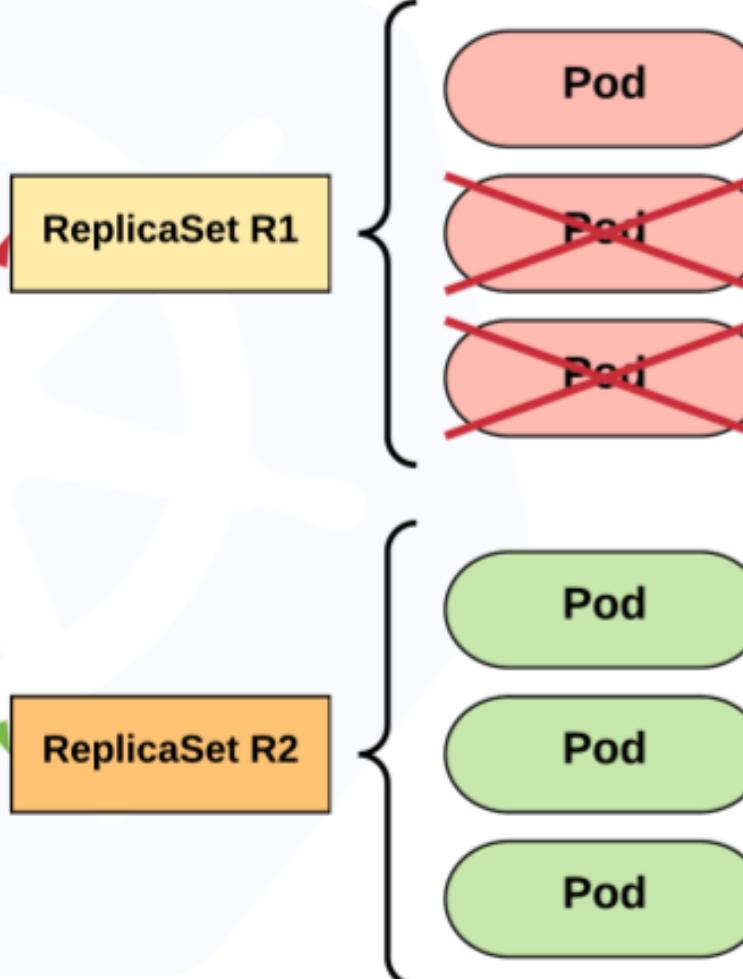
```
$ kubectl get replicaset  
NAME          DESIRED  CURRENT  READY   AGE  
mydep-54f7ff7d6d  3        3        3      10s  
mydep-676677fff  0        1        1      5h
```

```
$ kubectl get pods  
NAME          READY   STATUS    RESTARTS   AGE  
mydep-54f7ff7d6d-9gv1l  1/1     Running   0          7s  
mydep-54f7ff7d6d-cqvlq  1/1     Running   0          5s  
mydep-54f7ff7d6d-gccr6  1/1     Running   0          2s  
mydep-676677fff-9r2zn  1/1     Running   0          5h
```

Deployment  
Revision 2

ReplicaSet R1

ReplicaSet R2



# RollingUpdate Deployment

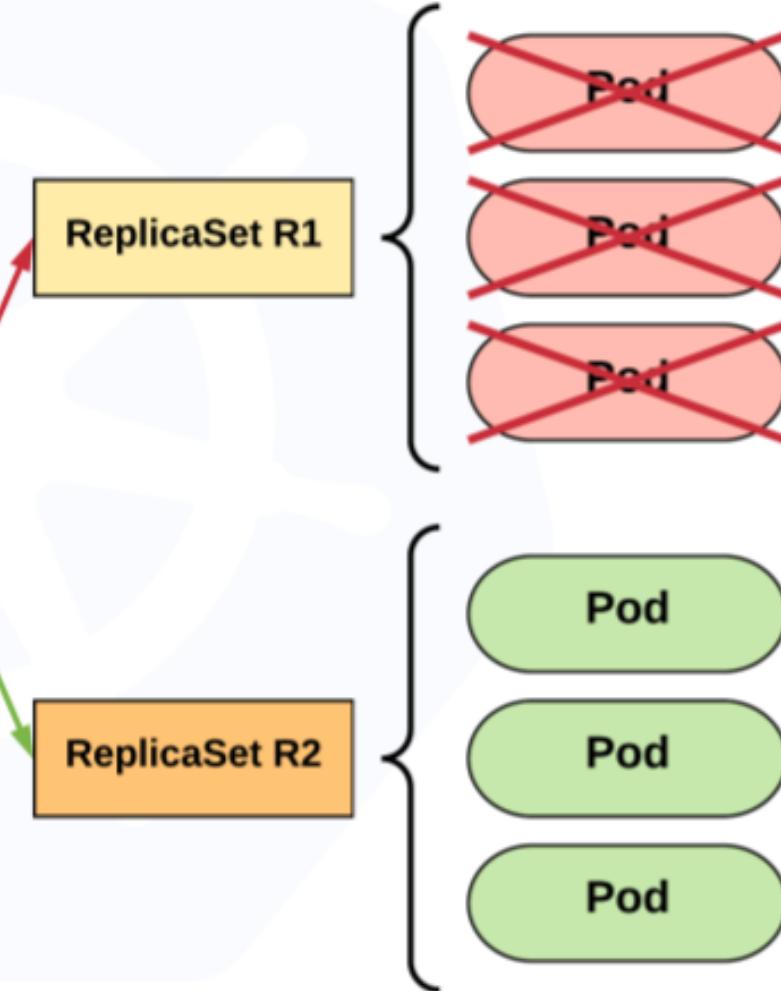
Phase out of old Pods managed by **maxSurge** and **maxUnavailable**.

R1 pod-template-hash:  
**676677fff**  
R2 pod-template-hash:  
**54f7ff7d6d**

Deployment  
Revision 2

```
$ kubectl get replicaset
NAME          DESIRED  CURRENT  READY   AGE
mydep-54f7ff7d6d  3        3        3      13s
mydep-676677fff  0        0        0      5h
```

```
$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
mydep-54f7ff7d6d-9gv1l  1/1     Running   0          10s
mydep-54f7ff7d6d-cqvlq  1/1     Running   0          8s
mydep-54f7ff7d6d-gccr6  1/1     Running   0          5s
```



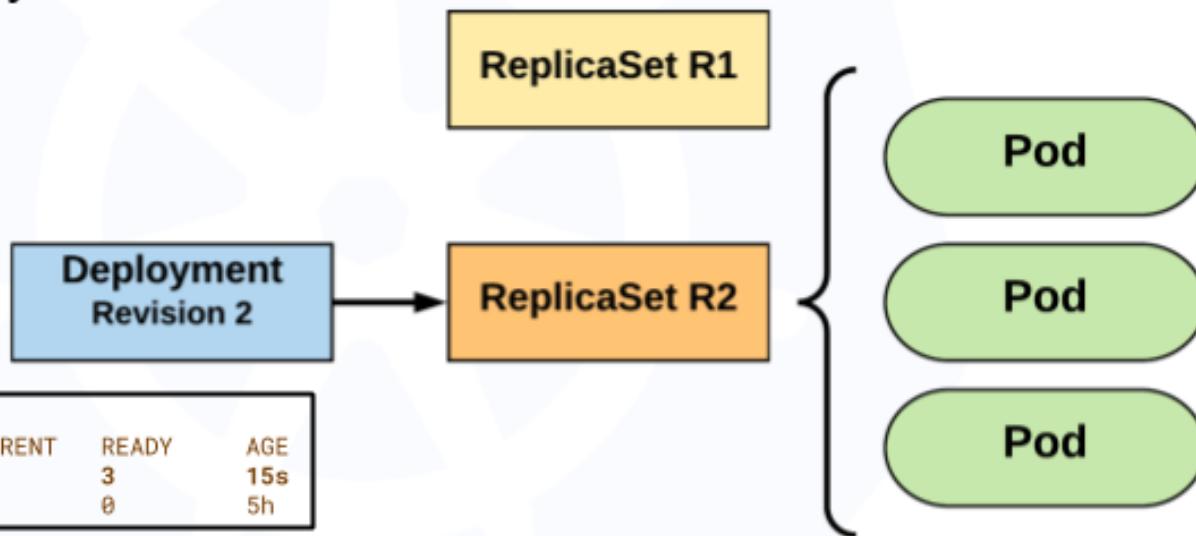
# RollingUpdate Deployment

Updated to new deployment revision completed.

```
R1 pod-template-hash:  
676677fff  
R2 pod-template-hash:  
54f7ff7d6d
```

```
$ kubectl get replicaset  
NAME          DESIRED  CURRENT  READY   AGE  
mydep-54f7ff7d6d  3        3        3      15s  
mydep-676677fff  0        0        0      5h
```

```
$ kubectl get pods  
NAME          READY   STATUS    RESTARTS   AGE  
mydep-54f7ff7d6d-9gv1l  1/1     Running   0          12s  
mydep-54f7ff7d6d-cqv1q  1/1     Running   0          10s  
mydep-54f7ff7d6d-gccr6  1/1     Running   0          7s
```



# Terminology - Deployments

- Create X Deployments units of nginx

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx
          ports:
            containerPort: 80
```

# Terminology - Deployments

- Use Deployment

```
kubectl scale deployment.v1.apps/nginx-deployment --replicas=10 --record=true
```

- Get details of your Deployment:

```
kubectl describe deployments
```

# Terminology - Services

---

- A **Service** is an **abstraction** which defines a
  - **Logical set** of Pods and a
  - **Policy** by which to access them (aka micro-service).
- Every service in K8S is actually a Load Balancer
- The set of Pods is accessed by a selector
- A **Service** in Kubernetes is a REST object
- A **Service** is responsible for enabling network access to a set of pods.
  - Pods can connect each other via direct network requests, while **Service** is used for exposing services to the "world"

# Terminology - Services

---

- Kubernetes assigns Service an IP address (aka “cluster IP”), which is used by the `Service` proxies
- `Services` “expose” the pods to the world
- `Services` serve as load balance layer
- Every node in a Kubernetes cluster runs a `kube-proxy` .  
kube-proxy is responsible for implementing a form of virtual IP for Services
- Services can define multiple ports and even static IPs
- When we deploy pods k8s they get the desired labels and the Service use selectors to match those labels (pods label).
- When deployment changes the service is auto updated and reflecting those changes (for example if we add more pods)

# Terminology - Services Ingress

---

- An ingress is a **host name and path mapping** to Service
- An Ingress is an object that allows access to your Kubernetes services from outside the Kubernetes cluster.
- Ingress is configured by creating a **collection of rules** that define which inbound connections reach which services.
- An ingress is usually but not only **LoadBalancer(s)** for exposing Node services.
- A basic deployment is usually build upon:
  - Deployment
  - Service
  - Ingress rules

# Terminology - Services

`kind: Service`

`apiVersion: v1`

`metadata:`

`name: hostname-service`

`spec:`

`type: NodePort`

`selector:`

`app: echo-hostname`

`ports:`

`- nodePort: 30163`

`port: 8080`

`targetPort: 80`

Make the service available  
to network requests from  
external clients

Forward requests to pods  
with label of this value

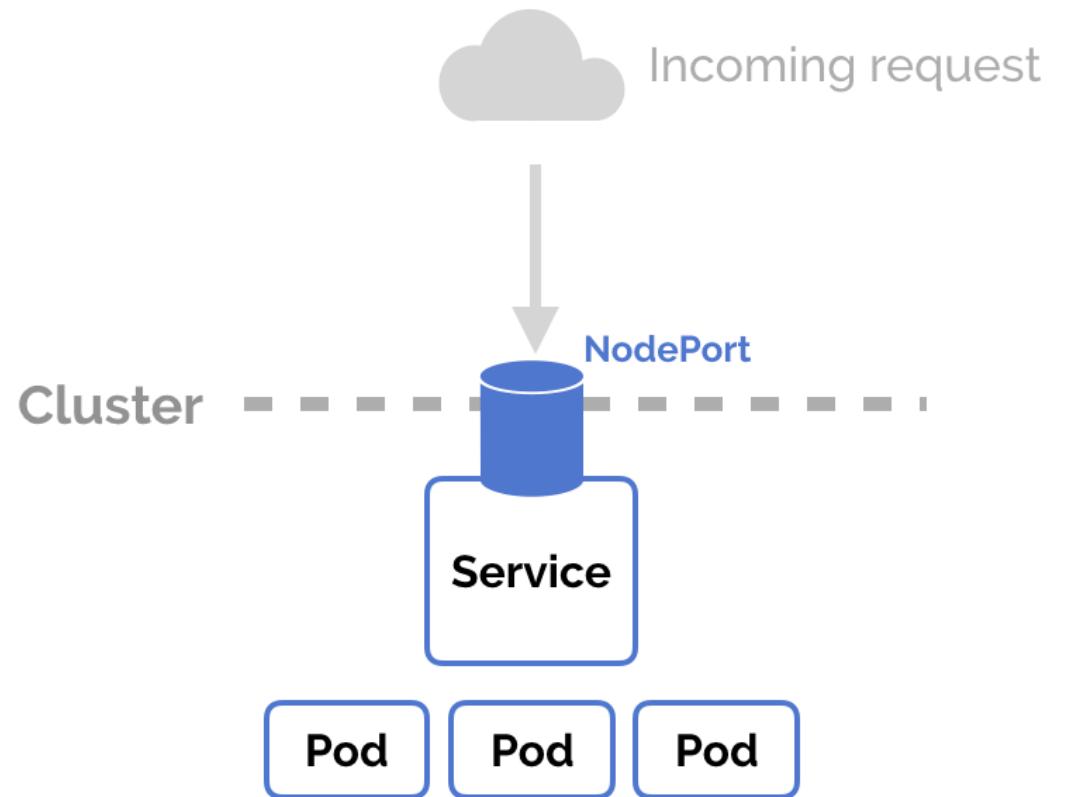
**nodePort**  
access service via this external port number

**port**  
port number exposed internally in cluster

**targetPort**  
port that containers are listening on

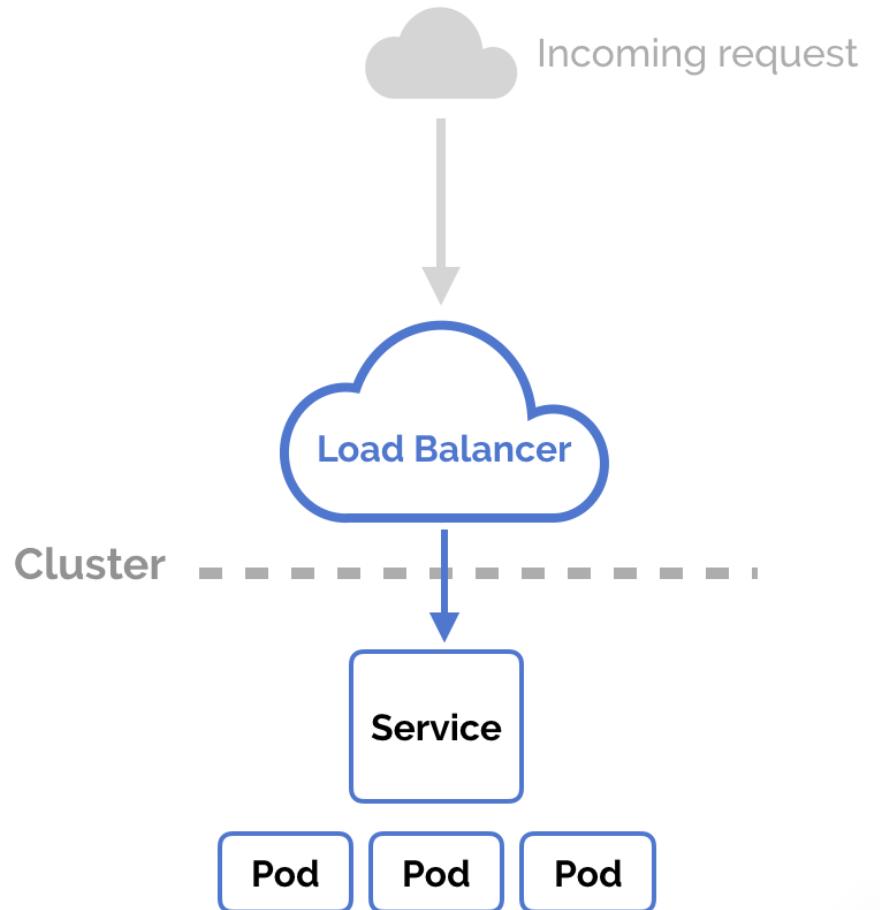
# Terminology - Services NodePort

- **NodePort** is a configuration setting you declare in a service's YAML.
- Set the service spec's type to **NodePort**. Then, Kubernetes will allocate a specific port on each Node to that service, and any request to your cluster on that port gets forwarded to the service.



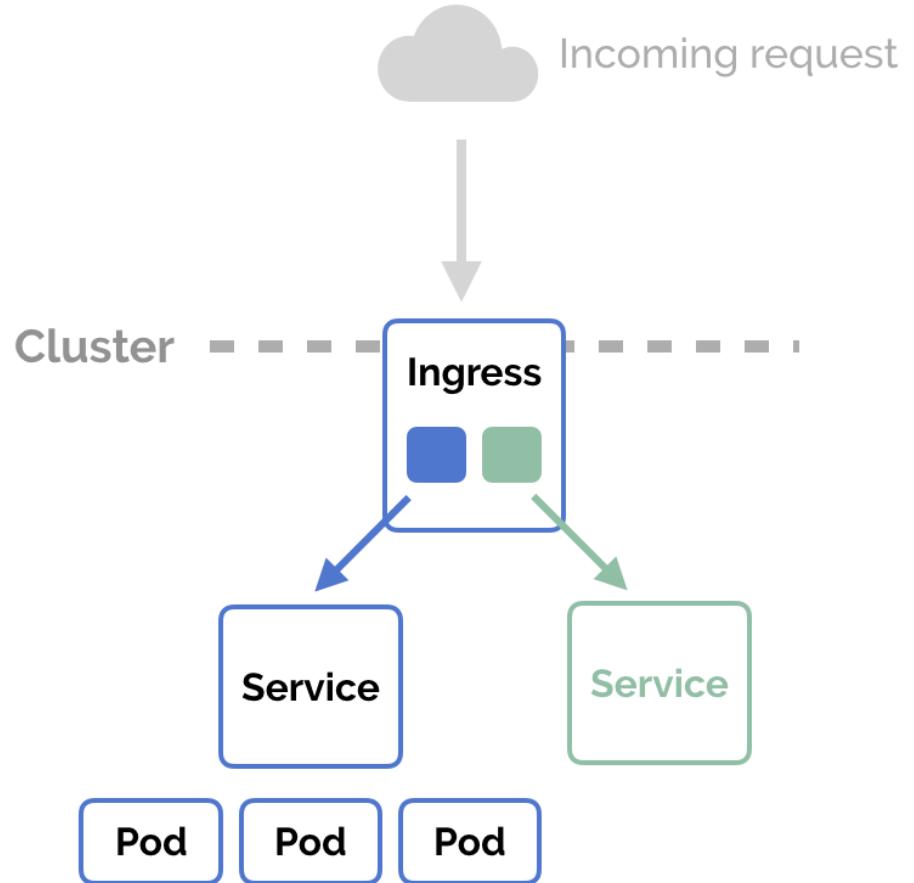
# Terminology - Services LoadBalancer

- **LoadBalancer** is usually a cloud provider LoadBalancer.
- Every time you want to expose a service to the outside world, you have to create a new LoadBalancer and get an IP address.



# Terminology - Services Ingress

**Ingress** Will be covered later on in the network section



# Terminology - Services yaml

```
kind: Service
apiVersion: v1
metadata:
  name: hostname-service
spec:
  # Expose the service on a static port on each node so that we can access the service from outside the cluster
  type: NodePort

  # When the node receives a request on the static port (30163) "select pods with the label 'app' set to 'echo-
  # hostname'"
  # and forward the request to one of them
  selector:
    app: echo-hostname

  # Three types of ports for a service
  #   - nodePort - a static port assigned on each the node
  #   - port - port exposed internally in the cluster
  #   - targetPort - the container port to send requests to
  ports:
    - nodePort: 30163
      port: 8080
      targetPort: 80
```

# Terminology - Services (ServiceTypes)

Kubernetes ServiceTypes allow you to specify what kind of Service you want.

Node types	Description
ClusterIP [default]	Exposes the Service on a cluster-internal IP, which makes the Service only reachable from within the cluster.
NodePort	Exposes the Service on each Node's IP at a static port (the NodePort).
LoadBalancer	Exposes the Service externally using a cloud provider's load balancer. NodePort and ClusterIP Services, to which the external load balancer routes, are automatically created.
ExternalName	Maps the Service to the contents of the externalName field (e.g. <a href="http://foo.bar.example.com">foo.bar.example.com</a> ), by returning a CNAME record

# Terminology - Volumes

---

- In Docker volumes are set of mapping paths from the host to the container and vice versa
- In k8s, volumes has an explicit lifetime - the same as the Pod that encloses it.  
For example: when Pod ceases to exist, the volume will cease to exist as well.
- k8s supports multiple types of volumes
  - My Favorite: **gitRepo** (deprecated)

# Terminology - Namespaces

---

- Kubernetes supports multiple virtual clusters deployed on the same physical cluster.  
(multiple separate sub groups)
- These virtual clusters are called namespaces .
- Namespaces are a way to divide cluster resources between multiple users (via resource quota).
- Namespaces provide a scope for names.  
Names of resources need to be unique within a namespace
- Namespaces can not be nested inside one another and each Kubernetes resource can only be in one namespace.

# Terminology - Namespaces

- Kubernetes starts with three initial namespaces:

Namespaces	Description
default	The default namespace for objects with no other namespace
kube-system	The namespace for objects created by the Kubernetes system
kube-public	This namespace is created automatically and is readable by all users (including those not authenticated).

# Terminology - ConfigMaps / Secrets

- **ConfigMaps** are used to pass and share key value pairs between pods
- **Secrets** are similar to config maps and store sensitive data

```
# Create ConfigMap explicitly
kubectl create configmap my-password --from-literal='password=123'

# Create ConfigMap from properties files
kubectl create configmap <config name> --from-file=<folder name>

# View config maps
kubectl get configmaps
```

# Terminology - ConfigMaps / Secrets

- Create ConfigMap from file

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: config-value01
  namespace: default
data:
  day: 'Sunday'
  year: '2019'
  raining: 'false'
```

```
# Create the config maps
kubectl create -f configmap.yaml
```

# Terminology - ConfigMaps / Secrets

- Read the config maps
- `env:` - Select a specific variable we wish to use

```
...
env:
- name: week-day
  valueFrom: <----->
    configMapKeyRef:
      name: <.....>
      key: <.....>
```

- `envFrom` - Select all the defined variables (in case we have multiple ones)

```
...
envFrom:
  ConfigMapRef:
    name: <.....>
```

# Terminology - ConfigMaps / Secrets

“ Note: Unlike ConfigMaps you must also specify the **type** of Secret you are creating. ,”

There are 3 types:

Type	Description
docker-registry	Credentials used to interact with a container registry.
generic	Equivalent to Opaque. Used for unstructured data.
tls	A TLS key pair (PEM Format) that accepts a cert (--cert) and key (--key).

# Terminology - ConfigMaps / Secrets

- Define secrets

```
# Create files needed for the secrets
echo -n 'admin' > ./username.txt
echo -n 'secret' > ./password.txt

# Create the secrets from the 2 separate files
kubectl create secret generic app-user-pass --from-file=./username.txt --from-file=./password.txt

# Verify that the secret is created
kubectl get secrets
```

# Terminology - Labels and selectors

- Labels are **key/value** pairs which attached and identify objects and resources in k8s.
- Every object in k8s should have a label
- Labels are used to organize and to select subsets of objects
- Each Key must be unique for a given object.
- Same objects will have the same label(s), for example instances in replica will have the same label
- Labels are used to mark releases, environments, versions ....

```
apiVersion: v1
kind: Pod
metadata:
  name: label-demo
labels:
  environment: production
  app: nginx
```

# Terminology - Selectors

- Selectors use labels to filter or select objects.

```
apiVersion: v1
kind: Pod
metadata:
  name: example
  labels:
    app: app1
    env: prod
spec:
  containers:
  - name: app1
    image: app1/app2
    ports:
    - containerPort: 80
nodeSelector:
  key: value
```

# Terminology - Labels and selectors

- Selectors types:
  - equality-based
  - set-based

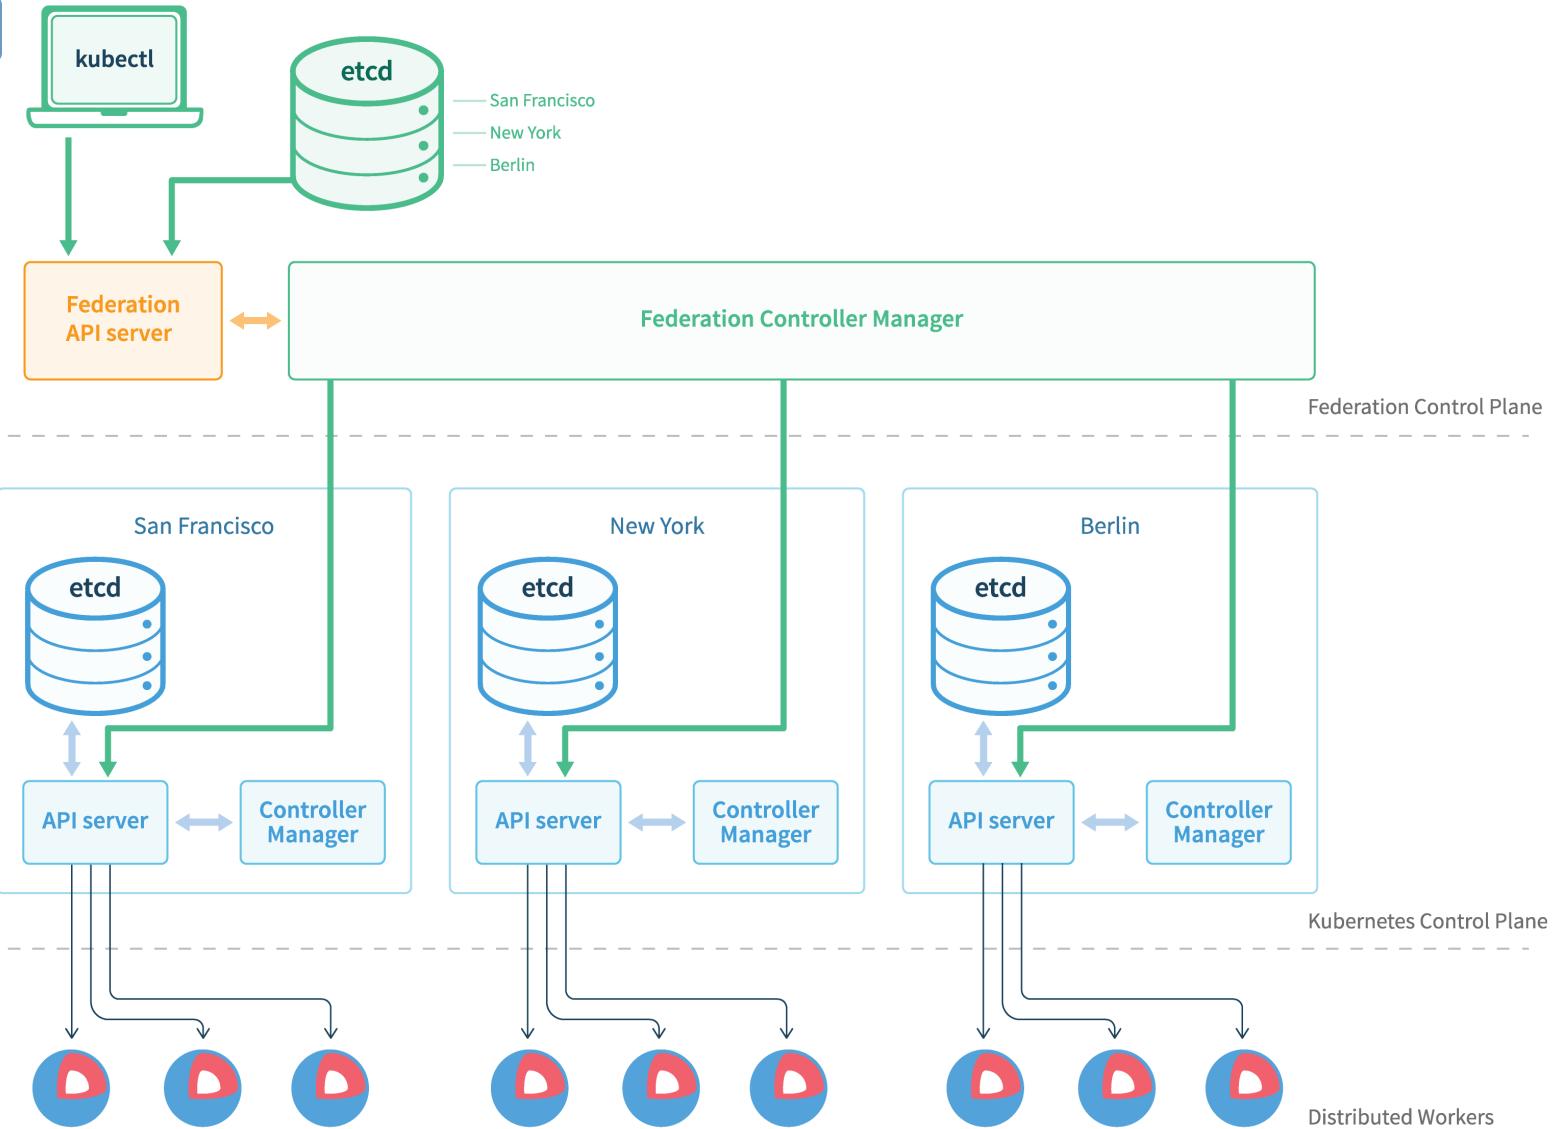
```
apiVersion: v1          selector:           selector:  
kind: Pod               matchLabels:        matchExpressions:  
metadata:              gpu: nvidia         - key: gpu  
  name: label-demo      |                  operator: in  
  labels:               |                  values: ["nvidia"]  
    environment: produc  
    app: nginx
```

# Terminology - etcd

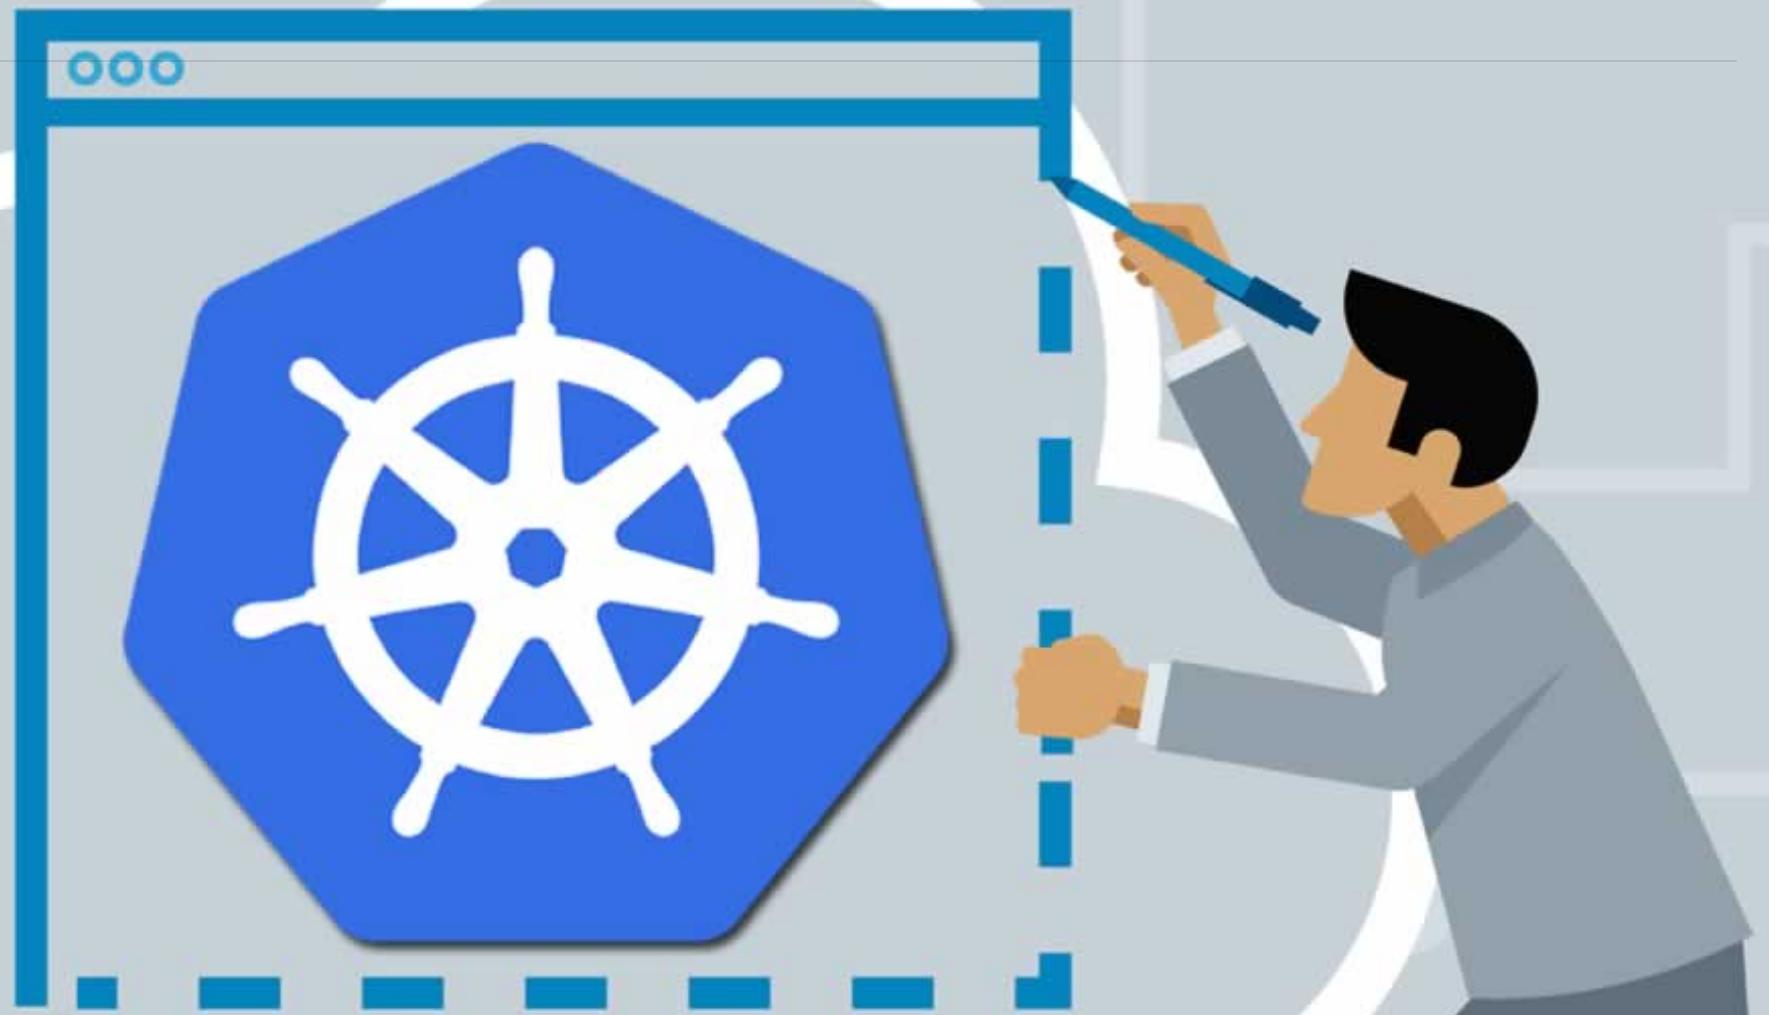
---

- **etcd** is a **consistent** and **highly-available** key value store used by Kubernetes
- **etcd** is string all cluster data.
- In production environment etcd should be replicated (there are several ways to implement it)

## etcd sample

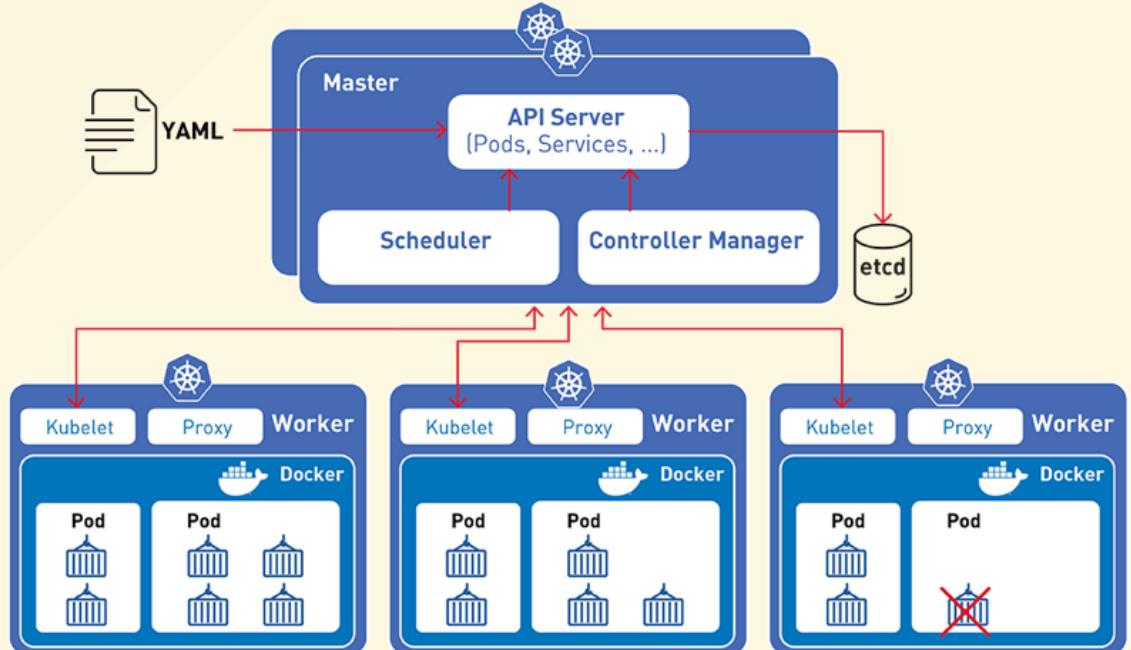


# K8S Intro



# How K8S works?

- K8s is build upon 4 main units
  - Pods
  - Services
  - Deployments / Replica sets
- User is interacting with the API
- K8S use yaml as the declarative language



# How K8S works?

---

- User describe the requirements in a yaml file ( [deployment](#) )
- The yaml is passed and processed by the API
- The K8S is responsible (auto managed) of creating, updating, scaling, healing and more
- The result is a cluster with the required objects
- The scheduler is responsible for the actual deployment
- In order to reach the pods (deployment units) we create [Service](#) .

# How K8S Deployment works?

---

- Deployment describe the cluster
- When we change the deployment, k8s initiate a rollback.
  - K8S notice the change
  - K8S gradually updates the number of pods and not all of them on the same time or we will loose our service
  - Pods has a grace period when they are deleted (default 30s)
  - K8S use the following methods to define healthy pods:
    - liveness check
    - readiness check

# How K8S Deployment works?

---

- K8S create new Pod
- K8S check for liveness

The first check is liveness check. Once the check pass - which mean that the deploy passed and the pod is live and running it continue to the next phase

- K8S now check for readiness check. Once this check pass the pod is added to the LB
- If K8S need to delete old pods it will give it a **termination grace period** (default 30s) before its being removed to finish processing current requests
  - The pod is moved to **termination state**
  - Once the termination period passed the pod is being deleted
- K8S repeat this process for every required pods

# How K8S Scheduler works? [kube-scheduler]

---

- The scheduler is the component which is responsible to deploy the required pod (on which machine)
- The scheduler is doing the following:
  - When we create a new pod the scheduler is "watching" for pods changes
  - It also monitor the state of all the machines (CPU, memory, storage etc)
- Once the scheduler notice a change it decide where to deploy the new pod based upon:
  - Predicates
  - Priorities

# How K8S Scheduler works? [kube-scheduler]

- **Predicates**

- Hard rule which must be accomplished like (OS, minimum memory etc).

Collection of

- System requirements (like memory)
  - User requirements (like - deploy only to VM with SSD drive)

# How K8S Scheduler works? [kube-scheduler]

- **Priorities**

Soft rules (recommendations) like:

- It would be nice if there will be a single pod on every VM
- It would be nice if you can deploy pods to different LB etc...
- Priorities usually include rules like spreading, taint (the machine is not 100% but still we allow to deploy to it)

For example: We prefer to deploy to a 100% percent healthy machine but if the only place left for deploy is a sick machine deploy anyway

- We can define scoring to each rule

# How K8S Scheduler works? [kube-scheduler]

---

- kube-scheduler selects a node for the pod in a 2-step operation:
  - Filtering
  - Scoring
  - `sort(filtering(nodes))` and choose the #1 node

# How Volumes and Storage works in K8S.

---

- A **Volume** is an object in each pod
- A **Volume** is associated with the pod and mapped to a path in the container
- There are different volumes in K8S
  - EmptyDir - Temporary volume, the folder is on the same machine as the pod
  - Mounted storage, locally or in the cloud which are shared between pods
  - Persistence volume (NFS, disks) etc. This volume lives even when the pod dies
- In order to "combine" the different volume types K8S use **Persistence Volume Claim**

# Persistence Volume Claim

---

- There are two ways PVs may be provisioned:
  - Statically
  - Dynamically.
- When Pod requires volume (storage) the K8S allocate (create) volume for each pod separately and mount them to the container.
- This allow more separately, flexibility between different storage types for example when we deploy to different cloud providers

# Secrets Management

---

- Secret is used for storing certificates, passwords etc
- A secret is a key value pair
  - The value can be string or file content
- Secret can be stored in volumes and we can declare specific secret volumes for serving sensitive data.
- Secret can also be an environment variable
- The content of the secret is stored in the etcd as plain text

# Secrets Management

---

- In latest versions K8S added **Key management store integration** for using encryption keys to encrypt sensitive data (ex. Azure key-vault)
- In case we use **Key management store integration** the data is encrypted before its being stored in the etcd

# Pods lifecycle



# Pods lifecycle & Phases

---

- A pod is the atomic unit of scheduling
- When we create the pod the first phase is `Pending`.  
Pending means that the Pod is "ready" (ex: yaml parsed, object are ready) but not yet deployed (not schedules)
- The scheduler "find" home for the pod as explained earlier `sort(filter(nodes))`
  - `Creating` - If required, the image is pulled from the registry
  - `Running` - Once the image is pulled the pod is running. During this phase in case we have error (failure) K8S try to restart the container several times.

# Pods lifecycle & Phases (Errors)

---

- In case of failure - when the retries doesn't work the pod state is changed to `CrashLoopBackOff`. K8S will try pause and might retry later.
- We can see the current states using `kubectl get pods`
- When pod stuck in `Pending` is usually due to lack of resources for deploying.
- When pod stuck in `Creating` is usually due to bad image
- `kubectl describe pods` will show the process of the pods

# Pods lifecycle & Phases (Advanced)

---

- Hooks
  - **PostStart** : Hook which is executed right after the container starts
  - **PreStop** : Hook which is executed before the container dies
- Init container
  - Init containers, run **before** the app containers are started.
  - Init containers are exactly like regular containers, except:
    - Init containers always run to completion.
    - Each init container must complete successfully before the next one starts.

# Pods (Advanced stuff - Probes)

---

- A Probe is a diagnostic performed periodically by the kubelet on a Container.

There are three types of probes for checking the pods:

- **ExecAction** : Executes a specified command inside the Container.  
The diagnostic is considered successful if the command exits with a status code of 0.
- **TCPSocketAction** : Performs a TCP check against the Container's IP address on a specified port.  
The diagnostic is considered successful if the port is open.
- **HTTPGetAction** : Performs an HTTP Get request.  
The diagnostic is considered successful if the response has a status code greater than or equal to 200 and less than 400.

# How Configuration management works in K8S.

---

- K8S use a **ConfigMaps**
- **ConfigMaps** are key-value pairs and can be created using yaml or properties file(s)
- **ConfigMaps** can be used as file or as environment variables
- **ConfigMaps** & **Secrets** are like but K8S handle them differently
- We define the configuration during the deployment
- Configuration change execute a rollup

# Monitoring & Alerting

---

- K8S expose many metrics for us (CPU, memory, disk etc)
- **Prometheus** is the leading tool for collecting data
  - We usually import the Prometheus and then add our monitoring metrics
  - Prometheus expose the metrics and we build the alerts based upon the metrics
- **Grafana** is used for visualization the metrics

# K8s API

---

- The K8S API is layer in the Master node
- The API server is the central management entity
- The API Server is the **only** component which is connected to the etcd

## kubectl

---

This is the api for end users for interacting with the k8c cluster

## kubelet

---

This component is part of each node and is used internally by the workers and the master

## kubeadm

---

Kubeadm is a tool for creating Kubernetes clusters.

# Rolling updates / Blue green deploy

---

- With blue/green deployments a **new version** (green) is deployed alongside the **existing version** (blue).
- Once the deploy is done then the ingress/router to the app is updated to switch to the new version (green).
- Sample code

# Networking (Basics)

---

- In K8S we use **Network Policies**
- **Network Policies** is a specification of how groups of pods are allowed to communicate with each other and other network endpoints.
- **NetworkPolicy** resources use labels to select pods and define rules which specify what traffic is allowed to the selected pods.
- The main policy types are
  - Ingress
  - Egress

# Networking (Basics)

---

- Ingress

Incoming traffic, **whitelist ingress rules**.

Each rule allows traffic which matches both the **from** and ports sections.

- Egress

Outgoing traffic, **whitelist egress rules**.

Each rule allows traffic which matches both the **to** and ports sections.

# Networking (Basics)

- By default, if no policies exist in a namespace, then **all ingress and egress** traffic is allowed to and from pods in the namespace.

## Network Policies

Policy	Description
allow-all	Allow all traffic to all pods in a namespace
default-deny	Block all traffic to all pods in a namespace

# Networking (Basics - Ingress)

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy

policyTypes:
- Ingress
- Egress
ingress:
- from:
- ipBlock:
    cidr: 172.17.0.0/16
    except:
    - 172.17.1.0/24
- namespaceSelector:
    matchLabels:
        project: myproject
- podSelector:
    matchLabels:
        role: frontend
ports:
- protocol: TCP
  port: 6379
```

# Networking (Basics - Egress)

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy

policyTypes:
- Ingress
- Egress
egress:
- to:
  - ipBlock:
    cidr: 10.0.0.0/24
ports:
- protocol: TCP
  port: 5978
```



# THE END.

